

# KaggleAZdataSet

November 5, 2020

```
[21]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
[2]: df = pd.read_csv('KaggleAZ/archive/A_Z Handwritten Data.csv')
```

```
[3]: df.head()
# zeroth column is the label
```

```
[3]:
```

	0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	...	0.639	0.640	0.641	\
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	

	0.642	0.643	0.644	0.645	0.646	0.647	0.648
0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0

[5 rows x 785 columns]

```
[4]: df.shape
```

```
[4]: (372450, 785)
```

```
[5]: X = df.iloc[:,1:]
y = df.iloc[:,0]
```

```
[6]: X
```

```
[6]:
```

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	0.10	...	0.639	0.640	\
0	0	0	0	0	0	0	0	0	0	0	...	0	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	0	

```

2      0      0      0      0      0      0      0      0      0      0      0 ...      0      0
3      0      0      0      0      0      0      0      0      0      0      0 ...      0      0
4      0      0      0      0      0      0      0      0      0      0      0 ...      0      0
...
372445  0      0      0      0      0      0      0      0      0      0      0 ...      0      0
372446  0      0      0      0      0      0      0      0      0      0      0 ...      0      0
372447  0      0      0      0      0      0      0      0      0      0      0 ...      0      0
372448  0      0      0      0      0      0      0      0      0      0      0 ...      0      0
372449  0      0      0      0      0      0      0      0      0      0      0 ...      0      0

```

```

          0.641  0.642  0.643  0.644  0.645  0.646  0.647  0.648
0          0      0      0      0      0      0      0      0
1          0      0      0      0      0      0      0      0
2          0      0      0      0      0      0      0      0
3          0      0      0      0      0      0      0      0
4          0      0      0      0      0      0      0      0
...
372445  0      0      0      0      0      0      0      0
372446  0      0      0      0      0      0      0      0
372447  0      0      0      0      0      0      0      0
372448  0      0      0      0      0      0      0      0
372449  0      0      0      0      0      0      0      0

```

[372450 rows x 784 columns]

```
[7]: y
```

```

[7]: 0      0
      1      0
      2      0
      3      0
      4      0
      ..
372445  25
372446  25
372447  25
372448  25
372449  25
Name: 0, Length: 372450, dtype: int64

```

```
[8]: X = X.astype('float32')
      y = y.astype('float32')
```

```
[9]: 28*28
```

```
[9]: 784
```

```
[10]: print(X.shape)
      print(y.shape)
```

```
(372450, 784)
(372450,)
```

```
[11]: X = X.values
      y = y.values
      X = X.reshape(X.shape[0],28,28,1)
      X/=255
```

```
[12]: input_shape = (28,28,1)
```

```
[13]: from sklearn.model_selection import train_test_split
```

```
[14]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
      ↪random_state=42)
```

```
[15]: import tensorflow as tf
      from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense,Conv2D,Dropout,Flatten, MaxPooling2D
```

```
[16]: model = Sequential()
      model.add(Conv2D(28, kernel_size=(3,3), input_shape=input_shape))
      model.add(MaxPooling2D(pool_size=(2, 2)))
      model.add(Flatten()) # Flattening the 2D arrays for fully connected layers
      model.add(Dense(128, activation=tf.nn.relu))
      model.add(Dropout(0.2))
      model.add(Dense(26,activation=tf.nn.softmax))
```

```
[17]: model.compile(optimizer='adam',
      loss='sparse_categorical_crossentropy',
      metrics=['accuracy'])
      model.fit(x=X_train,y=y_train, epochs=10)
```

```
Epoch 1/10
7799/7799 [=====] - 85s 11ms/step - loss: 0.2519 -
accuracy: 0.9278
Epoch 2/10
7799/7799 [=====] - 82s 10ms/step - loss: 0.1162 -
accuracy: 0.9663
Epoch 3/10
7799/7799 [=====] - 81s 10ms/step - loss: 0.0925 -
accuracy: 0.9722
Epoch 4/10
7799/7799 [=====] - 81s 10ms/step - loss: 0.0783 -
accuracy: 0.9765
```

```

Epoch 5/10
7799/7799 [=====] - 82s 10ms/step - loss: 0.0691 -
accuracy: 0.9787
Epoch 6/10
7799/7799 [=====] - 82s 10ms/step - loss: 0.0623 -
accuracy: 0.9804
Epoch 7/10
7799/7799 [=====] - 82s 10ms/step - loss: 0.0576 -
accuracy: 0.9818
Epoch 8/10
7799/7799 [=====] - 81s 10ms/step - loss: 0.0536 -
accuracy: 0.9831
Epoch 9/10
7799/7799 [=====] - 83s 11ms/step - loss: 0.0506 -
accuracy: 0.9842
Epoch 10/10
7799/7799 [=====] - 82s 11ms/step - loss: 0.0467 -
accuracy: 0.9850

```

[17]: <tensorflow.python.keras.callbacks.History at 0x7f8a1957fd60>

```
[18]: model.evaluate(X_test, y_test)
```

```

3841/3841 [=====] - 13s 3ms/step - loss: 0.0549 -
accuracy: 0.9863

```


[18]: [0.05487028881907463, 0.9863069653511047]

```
[22]: input_map = dict()
      c = 'a'
      for i in range(26):
          input_map[i] = c
          c = chr(ord(c)+1)
```

```
[23]: image_index = 2
      plt.imshow(X_test[image_index].reshape(28, 28), cmap='Greys')
      pred = model.predict(X_test[image_index].reshape(1, 28, 28, 1))
      print("Predicted Letter", input_map[pred.argmax()])
      print("Correct Label ", input_map[y_test[image_index]])
```


Predicted Letter c

Correct Label c




```
[24]: image_index = 20000
plt.imshow(X_test[image_index].reshape(28, 28),cmap='Greys')
pred = model.predict(X_test[image_index].reshape(1, 28, 28, 1))
print("Predicted Letter",input_map[pred.argmax()])
print("Correct Label ",input_map[y_test[image_index]])
```

Predicted Letter p  
Correct Label p




```
[25]: image_index = 89898
plt.imshow(X_test[image_index].reshape(28, 28),cmap='Greys')
pred = model.predict(X_test[image_index].reshape(1, 28, 28, 1))
print("Predicted Letter",input_map[pred.argmax()])
print("Correct Label ",input_map[y_test[image_index]])
```

Predicted Letter a  
Correct Label a



```
[26]: image_index = 1898
plt.imshow(X_test[image_index].reshape(28, 28),cmap='Greys')
pred = model.predict(X_test[image_index].reshape(1, 28, 28, 1))
print("Predicted Letter",input_map[pred.argmax()])
print("Correct Label ",input_map[y_test[image_index]])
```

Predicted Letter s  
Correct Label s



```
[27]: model.save('charRecognition.h5')
```

```
[ ]:
```

```
[28]: #Model 2
model = Sequential()
model.add(Conv2D(filters=32, kernel_size=(3,3),activation = 'relu',
    ↪,input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
model.add(Conv2D(filters=32, kernel_size=(3,3),activation = 'relu'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))

model.add(Flatten()) # Flattening the 2D arrays for fully connected layers

model.add(Dense(128, activation=tf.nn.relu))
model.add(Dense(26,activation=tf.nn.softmax))

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x=X_train,y=y_train, epochs=15)
```

```
Epoch 1/15
7799/7799 [=====] - 97s 12ms/step - loss: 0.1470 -
accuracy: 0.9579
```



```

Epoch 2/15
7799/7799 [=====] - 96s 12ms/step - loss: 0.0581 -
accuracy: 0.9837
Epoch 3/15
7799/7799 [=====] - 93s 12ms/step - loss: 0.0402 -
accuracy: 0.9881
Epoch 4/15
7799/7799 [=====] - 94s 12ms/step - loss: 0.0306 -
accuracy: 0.9905
Epoch 5/15
7799/7799 [=====] - 96s 12ms/step - loss: 0.0234 -
accuracy: 0.9927
Epoch 6/15
7799/7799 [=====] - 95s 12ms/step - loss: 0.0189 -
accuracy: 0.9939
Epoch 7/15
7799/7799 [=====] - 95s 12ms/step - loss: 0.0159 -
accuracy: 0.9949
Epoch 8/15
7799/7799 [=====] - 95s 12ms/step - loss: 0.0133 -
accuracy: 0.9956
Epoch 9/15
7799/7799 [=====] - 94s 12ms/step - loss: 0.0118 -
accuracy: 0.9961
Epoch 10/15
7799/7799 [=====] - 94s 12ms/step - loss: 0.0106 -
accuracy: 0.9964
Epoch 11/15
7799/7799 [=====] - 94s 12ms/step - loss: 0.0096 -
accuracy: 0.9969
Epoch 12/15
7799/7799 [=====] - 97s 12ms/step - loss: 0.0092 -
accuracy: 0.9970
Epoch 13/15
7799/7799 [=====] - 97s 12ms/step - loss: 0.0089 -
accuracy: 0.9971
Epoch 14/15
7799/7799 [=====] - 98s 13ms/step - loss: 0.0084 -
accuracy: 0.9975
Epoch 15/15
7799/7799 [=====] - 100s 13ms/step - loss: 0.0080 -
accuracy: 0.9975

```

[28]: <tensorflow.python.keras.callbacks.History at 0x7f8926b915b0>

```

[29]: model.evaluate(X_test, y_test)
      model.save('charRecognition2.h5')

```

3841/3841 [=====] - 14s 4ms/step - loss: 0.0548 -  
accuracy: 0.9920

[ ]: